

Università degli Studi di Padova – Department of Mathematics
Master's Degree in Data Science
Structural Bioinformatics, A.Y. 2022/23

Tanner Graves	2073559	tanner.graves@studenti.unipd.it
Marco Uderzo	2096998	marco.uderzo@studenti.unipd.it
Nour Alhousseini	2081230	nour.alhousseini@studenti.unipd.it
Hazeezat Adebimpe Adebayo	2090254	hazeezatadebimpe.adebayo@studenti.unipd.it

Classification of Contacts in Protein Structures

1. Training Set and Data Retrieval

The first step of the project was retrieving the data from each `.tsv` file from the `features_ring` folder (1,807 PDBs in total) and storing it into a single `DataFrame`.

The data contained in each `.tsv` file consists in a tab-separated file available for each protein, in which each row represents a contact in the protein and each column represents a feature about that contact. The last column is the target label, specifying the interaction type.

Column position	Column name	Column meaning	Type of column
1	pdb_id		
2	s_ch	chain	
3	s_resi	index	source residue identifier
4	s_ins	insertion code	
5	s_resn	name	
6	s_ss8	secondary structure 8 states (DSSP)	
7	s_rsa	relative solvent accessibility	
8	s_up	half sphere exposure up	
9	s_down	half sphere exposure down	
10	s_phi	phi angle	
11	s_psi	psi angle	
12	s_ss3	secondary structure 3 states (from angles)	source residue features
13	s_a1	Atchley feature 1	
14	s_a2	Atchley feature 2	
15	s_a3	Atchley feature 3	
16	s_a4	Atchley feature 4	
17	s_a5	Atchley feature 5	
18	t_ch	chain	
19	t_resi	index	target residue identifier
20	t_ins	insertion code	

21	t_resn	name	
22	t_ss8	secondary structure 8 states (DSSP)	
23	t_rsa	relative solvent accessibility	
24	t_up	half sphere exposure up	
25	t_down	half sphere exposure down	
26	t_phi	phi angle	
27	t_psi	psi angle	
28	t_ss3	secondary structure 3 states (from angles)	target residue features
29	t_a1	Atchley feature 1	
30	t_a2	Atchley feature 2	
31	t_a3	Atchley feature 3	
32	t_a4	Atchley feature 4	
33	t_a5	Atchley feature 5	
34	Interaction	interaction type	

Table 1.1: Training set description

Interaction Type	Count
HBOND	333,346
VDW	155,789
PIPISTACK	10,403
IONIC	9,068
SSBOND	866
PICATION	626
<i>Unclassified</i>	225,412

Table 1.2: Number of examples by bond type.

2. Data Preprocessing

The preprocessing pipeline starts with removing all samples where the label is unavailable (to explain how we are reinputting them). Then, the missing values of each feature are replaced using the mode of the feature itself. (Numerical features?)

A best subset selection is performed using Logistic Regression to determine which features are the most meaningful ones and the ones that are not influencing the decision much.

Scaling is then performed to standardize all the features to values between [0,1] to be then fed to the model.

The biggest criticality in the dataset is the heavy imbalance that is evident by looking at the number of contacts by interaction type. Notably, Hydrogen Bonds (HBOND) and Van der Waals contacts (VDW) are the most numerous, which makes them overrepresented in the training set. The remaining contact types, instead, are underrepresented. Training a model with such unbalanced datasets is sure to yield poor performance, especially when evaluating the model on new unseen data.

To mitigate this issue, a mixed approach of undersampling the most represented classes and oversampling the underrepresented ones is applied. For undersampling, an **InstanceHardnessThreshold** undersampler is used, which fits an estimator on the data and removes the most difficult data points to classify afterwards. For oversampling, SMOTE (*Synthetic Minority Oversampling TEchnique*) is used, which uses interpolation between samples to create new artificial data points.

It is to be noted, however, that altering the dataset, especially by oversampling, can yield overly optimistic results in the training performance, which don't necessarily transfer to as good performances at inference time on new data.

(Write about oversampling and undersampling parameters)

3. Model

3.1. Deep Neural Network

We decided to solve the classification of residue-residue contacts using Deep Learning, implementing a Deep Neural Network for multiclass classification. As for the library of choice, we used Keras, a very commonly used open-source deep learning library that acts as an interface for TensorFlow.

To perform multiclass classification using Neural Networks we need to encode the label of each sample into an identity vector, using a common practice called "One-Hot Encoding".

This ensures that each class is uniquely identified and independent of the others. It helps the neural network to better understand the categorical nature of the data and prevents any ordinal relationship assumptions between the classes.

One-Hot Encoding is also done to then set the number of neurons of the output layer to be equal to the number of classes in the dataset, in order to train each output neuron to determine the probability $P(C_i | \text{data})$, with $\sum P(C_i | \text{data}) = 1$. Therefore, the output layer represents the probability distribution of a contact being of a certain type.

(Write about implementation of L2 Regularization(?), Early Stopping...)

3.2. Model Details and Hyperparameters

Hyperparameters	
Batch Size	N
Weight Initialization	Xavier (Glorot Normal)
Loss Function	Categorical Cross-Entropy
Optimizer	Adam
Hidden Layers: Activation Function	ReLU
Output Layer: Activation Function	ReLU

Table 3.2.1: Model’s hyperparameters

Layer Type	Output Shape	Param #
Dense	64	
Dense	128	
Dense	N	
Dense	N	
Dense	N	
Dense	N	
# Total Parameters		174,918

Table 3.2.2: Model’s architecture.

4. Results

4.1 Performance

4.2 Issues

Implementing dataset balancing greatly improved model recall on minority classes. This would indicate that the model is effectively learning information about the classes. However, these classes still suffer from poor precision with resampling likely attributed to the relatively extremely low unique examples not allowing for a good representation of the class to be learned.

Through much experimentation it was found that more complex models with more parameters or lower dropout probabilities were better able to recall information about the minority classes, and simpler models would often not predict some of the minority classes at all, namely π -Cation interactions and SS-bonds.

The largest contributor to model inaccuracy was confusion between the two largest represented groups: hydrogen bonds and Van der Waals interactions.

5. Usage

To be written after we have the model running, and after having a model.py and main.py. main.py should also accept arguments to input data, retrain the network etc (implemented with an argument parser).

The full documentation of the software is available in the GitHub repository at the following link.